

SVIVVA

Vercel for AI APIs

Describe it. Get a live endpoint in 60 seconds.

Y COMBINATOR APPLICATION BRIEF

hello@svivva.com

WHAT WE BUILD

Svivva turns a plain-English prompt into a fully deployed, schema-validated AI API endpoint — no code required.

Developers paste a prompt describing what their API should do. Svivva generates the system prompt, enforces a strict JSON output schema, auto-creates 50–200 evaluation cases, runs them, and hands back a live public URL — in under 60 seconds. Every change is versioned with instant rollback if quality drops.

We are building the infrastructure layer that every team reaching for OpenAI's API actually needs: reliability, versioning, evaluation, and governance — without the glue code.

THE PROBLEM

Building production AI features today means:

- Writing fragile prompt strings buried in application code
- No versioning — one bad edit breaks production silently
- Hand-writing evaluation suites that never stay up to date
- Stitching together JSON validation, retry logic, and schema enforcement from scratch
- Zero visibility into latency, cost, or error rates across API calls

This costs mid-stage startups 2–4 engineering weeks per AI feature and produces brittle, untestable systems.

OUR SOLUTION

Svivva is a platform with five integrated layers:

% API Builder

Natural-language !' deployed endpoint in one flow. Schema enforcement via JSON Schema + AI auto-repair. OpenAPI 3.0 export and Python/Node SDK generation included.

% Evaluation Engine

Auto-generates 50–200 diverse test cases (edge, adversarial, happy-path) per endpoint. Runs on every deploy. Triggers instant rollback if pass-rate falls below threshold.

% Version Control

Every prompt or schema change creates a new immutable version. One-click rollback. Full diff history.

% Neural AI Pipeline

Prompt Optimizer, Schema Enhancer, Quality Gate (confidence scoring), Training Augmentation (5 synthetic-data strategies), Anomaly Detector, Chaos Mode stress testing, Prompt Breeding (genetic evolution), API Autopsy (failure forensics).

% Live Public Runtime

Every deployed API gets a CORS-enabled public endpoint at `/api/run/[slug]`. Returns structured JSON, latency, and a "powered by Svivva" attribution. Shareable API card page for each endpoint.

COMPLETE FEATURE SET

Core API Platform

- Prompt-to-endpoint builder — system prompt + output schema from plain English
- JSON Schema enforcement with AJV validation and AI-powered auto-repair on malformed output
- OpenAPI 3.0 spec export (YAML/JSON)
- Auto-generated Python and Node.js client SDKs
- Secure API key management with rotation
- Real-time usage analytics: latency P50/P95, cost per call, error rate
- Configurable usage alerts via email and Slack webhooks
- One-click deploy !' live public /api/run/[slug] endpoint with CORS
- Shareable public API card page (prompt, schema, cURL snippet, embed badge)

Evaluation & Quality

- Auto-generation of 50–200 evaluation cases per endpoint (edge, adversarial, happy-path)
- AI-graded test examples with scoring rubrics
- Automated evaluation runner on every deploy
- Instant rollback triggered when pass-rate drops below configurable threshold
- Quality Gate: confidence scoring on every live API response
- Chaos Mode: adversarial stress testing with resilience score report

Version Control & Collaboration

- Immutable versioning — every prompt/schema change creates a new version
- One-click rollback to any prior version
- A/B testing: traffic-split between two prompt versions with performance comparison
- Team management: role-based access (owner, admin, member, viewer)
- Project-level and team-level API key scoping

Neural AI Pipeline

- Prompt Optimizer: AI analysis and rewrite of system prompts for higher quality
- Schema Enhancer: AI suggestions to tighten or expand output schemas
- Training Augmentation: synthetic data generation using 5 distinct strategies
- Anomaly Detector: pattern recognition across logs to surface failure clusters
- API Autopsy: forensic root-cause analysis of individual API failures
- Prompt Breeding: genetic algorithm that cross-breeds two prompt versions
- Fine-tuning pipeline: export training examples for custom model training
- Cost Optimizer: model recommendation engine with budget capping

API Marketplace

- Publish any API to a public marketplace for discovery and monetisation
- Per-use or subscription pricing on marketplace listings
- Ratings, forks, and usage analytics per listing

Idea Engine

- AI-powered market scanning !' gap analysis !' idea generation pipeline
- Generates 6 novel ideas per session with novelty, revenue, and feasibility scores
- Industry filtering and custom context input
- Full session history for revisiting past discoveries

Svivva Seeds — Multi-App Generation Engine

- Upload one structured PDF !' generate multiple complete production apps in parallel
- Generates frontend, backend, database schema, auth, and deployment config per app
- Extracts engineering docs (API docs, architecture, deployment guides) automatically
- Extracts marketing content (landing pages, pitch decks, app store descriptions)
- Multi-app prompt editing: select N seeds, apply one AI prompt to all specs in parallel
- Auto Marketing Engine: 3 SEO landing pages per app published to live root-level slugs
- Sitemap auto-pinged to Google on every generation run
- Available on Pro and Enterprise tiers

Hardware Builder (Physical Products)

- Parallel Digital + Physical mode — same platform for software APIs and hardware product specs
- AI-generated hardware product specifications from prompts
- Component list, BOM estimation, and 3D model scaffolding
- Separate physical-mode sidebar and workflow

SEO & Growth Infrastructure

- Programmatic SEO: server-rendered tool pages at /tools/[slug]
- Blog with JSON-LD structured data, category filters, related posts, CTA sections
- AI landing page generator from keywords (saved to DB, published at root slugs)
- AI article generator from keywords
- GA4 + Google Ads conversion tracking built in
- Dynamic sitemap.xml and robots.txt
- Orbit admin launchpad: 12-step Svivva traffic plan + 5-step Mini Apps plan

Svivva Play (Music Production Tool)

- Browser-based music production tool with hardware-sampler UI
- Deterministic Neo-Soul chord engine (Glasper cyclic-3rds, Ivan Lins sus13 voicings)
- Correct MIDI note numbers, separate bass stem generation
- Audio file import with iOS audio pre-warm on interaction
- Coming-soon overlay; accessible on Pro/Enterprise tiers

BUSINESS MODEL

Free \$0/mo

Up to 3 API endpoints, 1,000 API calls/mo, 30-day version history

Pro \$49/mo

Unlimited endpoints, 100k API calls/mo, evaluations, version control, Svivva Seeds, Svivva Play, A/B testing

Enterprise Custom

SSO, SLA, dedicated infra, unlimited everything, on-prem option, team management, priority support

Marketplace Rev-share

20% platform fee on all marketplace API transactions; zero upfront for sellers

WHY NOW

Every company is racing to ship AI features. OpenAI, Anthropic, and Google now offer capable models as commodity infrastructure — the bottleneck has shifted to tooling: reliability, evaluation, versioning, and governance. The market for AI dev-tools is nascent and winner-take-most. Svivva is the first platform to bundle all of these layers into a single product aimed at teams that want to ship AI APIs as confidently as they ship REST APIs.

TECHNICAL STACK

Frontend: Next.js 16 App Router, TypeScript, Tailwind CSS, shadcn/ui, Three.js

Backend: Next.js API Routes, Node.js

Database: PostgreSQL + Drizzle ORM (48 tables)

AI: OpenAI GPT-4o via Replit AI Integrations

Auth: Replit OIDC (JWT session)

Payments: Stripe

Validation: Zod + AJV (JSON Schema)

Deployment: Replit cloud hosting; public endpoints at `/api/run/[slug]`

CONTACT

hello@svivva.com